

Lab 4. Deep Dive into Kafka Producers & Consumers



We will cover the following topics in this lab:

- The Kafka Producer API and its uses
- Additional configuration for producers
- An example of a producer
- Kafka consumer internals
- Kafka consumer APIs
- Java Kafka consumer example

Lab Solution

Complete solution for this lab is available in the following directory:

```
~/kafka-advanced/labs/Lab04
```

Kafka Producer APIs

Creating a Kafka producer involves the following steps:

1. Required configuration.
2. Creating a producer object.
3. Setting up a producer record.
4. Creating a custom partition if required.
5. Additional configuration.

Required configuration: In most applications, we first start with creating the initial configuration without which we cannot run the application. The following are three mandatory configuration parameters:

- `bootstrap.servers` : This contains a list of Kafka brokers addresses. The address is specified in terms of `hostname:port` . We can specify one or more broker detail, but we recommend that you provide at least two so that if one broker goes down, producer can use the other one.

Note

It is not necessary to specify all brokers as the Kafka producer queries this configured broker for information about other brokers. In older versions of Kafka, this property was `metadata.broker.list` , where we used to specify a list of brokers `host:port` .

- `key.serializer` : The message is sent to Kafka brokers in the form of a key-value pair. Brokers expect this key-value to be in byte arrays. So we need to tell producer which serializer class is to be used to convert this key-value object to a byte array. This property is set to tell the producer which class to use to serialize the key of the message.

Kafka provides us with three inbuilt serializer classes: `ByteArraySerializer` , `StringSerializer` , and `IntegerSerializer` . All these classes are present in the `org.apache.kafka.common.serialization` package and implement the serializer interface.

- `value.serializer` : This is similar to the `key.serializer` property, but this property tells the producer which class to use in order to serialize the value. You can implement your own serialize class and assign to this property.

Let's see how we do it in a programming context.

Here is how Java works for Producer APIs:

```
Properties producerProps = new Properties();
producerProps.put("bootstrap.servers", "broker1:port,broker2:port");
producerProps.put("key.serializer",
    "org.apache.kafka.common.serialization.StringSerializer");
producerProps.put("value.serializer",
    "org.apache.kafka.common.serialization.StringSerializer");
KafkaProducer<String, String> producer = new KafkaProducer<String,String>
(producerProps);
```

Producer object and ProducerRecord object

Producer accepts the `ProducerRecord` object to send records to the `.ProducerRecord` topic. It contains a topic name, partition number, `timestamp`, key, and value. Partition number, `timestamp`, and key are optional parameters, but the topic to which data will be sent and value that contains the data is mandatory.

- If the partition number is specified, then the specified partition will be used when sending the record
- If the partition is not specified but a key is specified, a partition will be chosen using a hash of the key
- If both key and partition are not specified, a partition will be assigned in a round-robin fashion

Here is the `producerRecord` in Java:

```
ProducerRecord producerRecord = new ProducerRecord<String, String>(topicName, data);
Future<RecordMetadata> recordMetadata = producer.send(producerRecord);
```

Here is an example of `producerRecord` in Scala:

```
val producerRecord = new ProducerRecord<String, String>(topicName, data);
val recordMetadata = producer.send(producerRecord);
```

We have different constructors available for `ProducerRecord`:

- Here is the first constructor for `producerRecord`:

```
ProducerRecord(String topicName, Integer numberOfpartition, K key, V value)
```

- The second constructor goes something like this:

```
ProducerRecord(String topicName, Integer numberOfpartition, Long timestamp, K key, V value)
```

- The third constructor is as follows:

```
ProducerRecord(String topicName, K key, V value)
```

- The final constructor of our discussion is as follows:

```
ProducerRecord(String topicName, V value)
```

Each record also has a `timestamp` associated with it. If we do not mention a `timestamp`, the producer will stamp the record with its current time. The `timestamp` eventually used by Kafka depends on the `timestamp`

type configured for the particular topic:

- **[CreateTime]:** The `timestamp` of `ProducerRecord` will be used to append a `timestamp` to the data
- **[LogAppendTime]:** The Kafka broker will overwrite the `timestamp` of `ProducerRecord` to the message and add a new `timestamp` when the message is appended to the log

Once data is sent using the `send()` method, the broker persists that message to the partition log and returns `RecordMetadata`, which contains metadata of the server response for the record, which includes `offset`, `checksum`, `timestamp`, `topic`, `serializedKeySize`, and so on. We previously discussed common messaging publishing patterns. The sending of messages can be either synchronous or asynchronous.

[Synchronous messaging]: Producer sends a message and waits for brokers to reply. The Kafka broker either sends an error or `RecordMetadata`. We can deal with errors depending on their type. This kind of messaging will reduce throughput and latency as the producer will wait for the response to send the next message.

Generally, Kafka retries sending the message in case certain connection errors occur. However, errors related to serialization, message, and so on have to be handled by the application, and in such cases, Kafka does not try to resend the message and throws an exception immediately.

Java:

```
ProducerRecord producerRecord = new ProducerRecord<String, String>(topicName, data);

Object recordMetadata = producer.send(producerRecord).get();
```

[Asynchronous messaging]: Sometimes, we have a scenario where we do not want to deal with responses immediately or we do not care about losing a few messages and we want to deal with it after some time.

Kafka provides us with the callback interface that helps in dealing with message reply, irrespective of error or successful. `send()` can accept an object that implements the callback interface.

```
send(ProducerRecord<K,V> record,Callbackcallback)
```

The callback interface contains the `onCompletion` method, which we need to override. Let's look at the following example:

Here is the example in Java:

```
public class ProducerCallback implements Callback {
    public void onCompletion(RecordMetadata recordMetadata, Exception ex) {
        if(ex!=null){
            //deal with exception here
        }
        else{
            //deal with RecordMetadata here
        }
    }
}
```

Once we have the `Callback` class implemented, we can simply use it in the `send` method as follows:

```
val callBackObject = producer.send(producerRecord,new ProducerCallback());
```

If Kafka has thrown an exception for the message, we will not have a null exception object. We can also deal with successful and error messages accordingly in `onCompletion()`.

Custom partition

Remember that we talked about key serializer and value serializer as well as partitions used in Kafka producer. As of now, we have just used the default partitioner and inbuilt serializer. Let's see how we can create a custom partitioner.

Kafka generally selects a partition based on the hash value of the key specified in messages. If the key is not specified/null, it will distribute the message in a round-robin fashion. However, sometimes you may want to have your own partition logic so that records with the same partition key go to the same partition on the broker. We will see some best practices for partitions later in this lab. Kafka provides you with an API to implement your own partition.

In most cases, a hash-based default partition may suffice, but for some scenarios where a percentage of data for one key is very large, we may be required to allocate a separate partition for that key. This means that if key K has 30 percent of total data, it will be allocated to partition N so that no other key will be assigned to partition N and we will not run out of space or slow down. There can be other use cases as well where you may want to write `Custom Partition`. Kafka provides the `PartitionerInterface`, which helps us create our own partition.

Here is an example in Java:

```
public class CustomPartition implements Partitioner {
    public int partition(String topicName, Object key, byte[] keyBytes, Object value,
byte[] valueByte, Cluster cluster) {
        List<PartitionInfo> partitions = cluster.partitionsForTopic(topicName);

        int numPartitions = partitions.size();
        //Todo: Partition logic here
        return 0;
    }

    public void close() {

    }

    public void configure(Map<String, ?> map) {

    }
}
```

Additional producer configuration

There are other optional configuration properties available for Kafka producer that can play an important role in performance, memory, reliability, and so on:

- `buffer.memory` : This is the amount of memory that producer can use to buffer a message that is waiting to be sent to the Kafka server. In simple terms, it is the total memory that is available to the Java producer to collect unsent messages. When this limit is reached, the producer will block the messages for `max.block.ms` before raising an exception. If your batch size is more, allocate more memory to the producer buffer.

Additionally, to avoid keeping records queued indefinitely, you can set a timeout using `request.timeout.ms`. If this timeout expires before a message can be successfully sent, then it will be removed from the queue and an exception will be thrown.

- `acks` : This configuration helps in configuring when producer will receive acknowledgment from the leader before considering that the message is committed successfully:
 - `acks=0` : Producer will not wait for any acknowledgment from the server. Producer will not know if the message is lost at any point in time and is not committed by the leader broker. Note that no retry will happen in this case and the message will be completely lost. This can be used when you want to achieve very high throughput and when you don't care about potential message loss.
 - `acks=1` : Producer will receive an acknowledgment as soon as the leader has written the message to its local log. If the leader fails to write the message to its log, producer will retry sending the data according to the retry policy set and avoid potential loss of messages. However, we can still have message loss in a scenario where the leader acknowledges to producer but does not replicate the message to the other broker before it goes down.
 - `acks=all` : Producer will only receive acknowledgment when the leader has received acknowledgment for all the replicas successfully. This is a safe setting where we cannot lose data if the replica number is sufficient to avoid such failures. Remember, throughput will be lesser than the first two settings.
- `batch.size` : This setting allows the producer to batch the messages based on the partition up to the configured amount of size. When the batch reaches the limit, all messages in the batch will be sent. However, it's not necessary that producer wait for the batch to be full. It sends the batch after a specific time interval without worrying about the number of messages in the batch.
- `linger.ms` : This represents an amount of time that a producer should wait for additional messages before sending a current batch to the broker. Kafka producer waits for the batch to be full or the configured `linger.ms` time; if any condition is met, it will send the batch to brokers. Producer will wait till the configured amount of time in milliseconds for any additional messages to get added to the current batch.
- `compression.type` : By default, producer sends uncompressed messages to brokers. When sending a single message, it will not make that much sense, but when we use batches, it's good to use compression to avoid network overhead and increase throughput. The available compressions are GZIP, Snappy, or LZ4. Remember that more batching would lead to better compression.
- `retires` : If message sending fails, this represents the number of times producer will retry sending messages before it throws an exception. It is irrespective of reseeding a message after receiving an exception.
- `max.in.flight.requests.per.connection` : This is the number of messages producer can send to brokers without waiting for a response. If you do not care about the order of the messages, then setting its value to more than 1 will increase throughput. However, ordering may change if you set it to more than 1 with retry enabled.
- `partitioner.class` : If you want to use a custom partitioner for your producer, then this configuration allows you to set the partitioner class, which implements the partitioner interface.
- `timeout.ms` : This is the amount of time a leader will wait for its followers to acknowledge the message before sending an error to producer. This setting will only help when `acks` is set to all.

Java Kafka producer example

We have covered different configurations and APIs in previous sections. Let's start coding one simple Java producer, which will help you create your own Kafka producer.

[Prerequisite]

- IDE: IntelliJ Idea.
- Build tool: Maven, Gradle, or others. We have used Maven to build our project.
- `Pom.xml` : Add Kafka dependency to the `pom` file:

```
<dependency>
  <groupId>org.apache.kafka</groupId>
  <artifactId>kafka_2.11</artifactId>
  <version>0.10.0.0</version>
</dependency>
```

Java:

```
import java.util.Properties;
import java.util.concurrent.Future;
import org.apache.kafka.clients.producer.KafkaProducer;
import org.apache.kafka.clients.producer.ProducerRecord;
import org.apache.kafka.clients.producer.RecordMetadata;

public class DemoProducer {

    public static void main(final String[] args) {
        Properties producerProps = new Properties();
        producerProps.put("bootstrap.servers", "localhost:9092");
        producerProps.put("key.serializer",
"org.apache.kafka.common.serialization.StringSerializer");
        producerProps.put("value.serializer",
"org.apache.kafka.common.serialization.StringSerializer");
        producerProps.put("acks", "all");
        producerProps.put("retries", 1);
        producerProps.put("batch.size", 20000);
        producerProps.put("linger.ms", 1);
        producerProps.put("buffer.memory", 24568545);
        KafkaProducer<String, String> producer = new KafkaProducer<String, String>
(producerProps);

        for (int i = 0; i < 2000; i++) {
            ProducerRecord data = new ProducerRecord<String, String>("test1", "Hello
this is record " + i);
            Future<RecordMetadata> recordMetadata = producer.send(data);
        }
        producer.close();
    }
}
```

The preceding example is a simple Java producer where we are producing string data without a key. We have also hardcoded the topic name, which probably can be read through configuration file or as an command line input. To understand producer, we have kept it simple. However, we will see good examples in upcoming labs where we will follow good coding practice.

Consumer configuration

Let's see how we set and configure consumer in the real programming world.

Java:

```
Properties consumerProperties = new Properties();
consumerProperties.put("bootstrap.servers", "localhost:9092");
```

```
consumerProperties.put("group.id", "Demo");
consumerProperties.put("key.deserializer", "org.apache.kafka.common.serialization.StringDeserializer");
String> consumer = new KafkaConsumer<String, String>(consumerProperties);
```

Committing and polling

Polling is fetching data from the Kafka topic. Kafka returns the messages that have not yet been read by consumer. How does Kafka know that consumer hasn't read the messages yet?

Java:

```
while (true) {
    ConsumerRecords<String, String> records = consumer.poll(2);
    for (ConsumerRecord<String, String> record : records)
        System.out.printf("offset = %d, key = %s, value = %sn",
            record.offset(), record.key(), record.value());
    try {
        consumer.commitSync();
    } catch (CommitFailedException ex) {
        //Logger or code to handle failed commit
    }
}
```

- **[Asynchronous commit]:** The problem with synchronous commit is that unless we receive an acknowledgment for a commit offset request from the Kafka server, consumer will be blocked. This will cost low throughput. It can be done by making commit happen asynchronously. However, there is a problem in asynchronous commit--it may lead to duplicate message processing in a few cases where the order of the commit offset changes. For example, offset of message 10 got committed before offset of message 5. In this case, Kafka will again serve message 5-10 to consumer as the latest offset 10 is overridden by 5.

Java:

```
while (true) {
    ConsumerRecords<String, String> records = consumer.poll(2);
    for (ConsumerRecord<String, String> record : records)
        System.out.printf("offset = %d, key = %s, value = %sn",
            record.offset(), record.key(), record.value());
    consumer.commitAsync(new OffsetCommitCallback() {
        public void onComplete(Map<TopicPartition, OffsetAndMetadata> map, Exception e) {

        }
    });
}
```

You have learned about synchronous and asynchronous calls. However, the best practice is to use a combination of both. Asynchronous should be used after every poll call and synchronous should be used for behaviors such as the triggering of the rebalancer, closing consumer due to some condition, and so on.

Kafka also provides you with an API to commit a specific offset.

Java Kafka consumer

The following program is a simple Java consumer which consumes data from topic test. Please make sure data is already available in the mentioned topic otherwise no record will be consumed.

```
import org.apache.kafka.clients.consumer.*;
import org.apache.kafka.common.TopicPartition;
import org.apache.log4j.Logger;

import java.util.*;

public class DemoConsumer {
    private static final Logger log = Logger.getLogger(DemoConsumer.class);

    public static void main(String[] args) throws Exception {

        String topic = "test1";
        List<String> topicList = new ArrayList<>();
        topicList.add(topic);
        Properties consumerProperties = new Properties();
        consumerProperties.put("bootstrap.servers", "localhost:9092");
        consumerProperties.put("group.id", "Demo_Group");
        consumerProperties.put("key.deserializer",
            "org.apache.kafka.common.serialization.StringDeserializer");
        consumerProperties.put("value.deserializer",
            "org.apache.kafka.common.serialization.StringDeserializer");

        consumerProperties.put("enable.auto.commit", "true");
        consumerProperties.put("auto.commit.interval.ms", "1000");
        consumerProperties.put("session.timeout.ms", "30000");

        KafkaConsumer<String, String> demoKafkaConsumer = new KafkaConsumer<String,
String>(consumerProperties);

        demoKafkaConsumer.subscribe(topicList);
        log.info("Subscribed to topic " + topic);
        int i = 0;
        try {
            while (true) {
                ConsumerRecords<String, String> records = demoKafkaConsumer.poll(500);
                for (ConsumerRecord<String, String> record : records)
                    log.info("offset = " + record.offset() + "key =" + record.key() +
"value =" + record.value());

                //TODO : Do processing for data here
                demoKafkaConsumer.commitAsync(new OffsetCommitCallback() {
                    public void onComplete(Map<TopicPartition, OffsetAndMetadata> map,
Exception e) {

                    }
                });
            }
        } catch (Exception ex) {
```



```

        //TODO : Log Exception Here
    } finally {
        try {
            demoKafkaConsumer.commitSync();

        } finally {
            demoKafkaConsumer.close();
        }
    }
}
}
}

```

Rebalance listeners

We discussed earlier that in case of addition or removal of consumer to the consumer group, Kafka triggers the rebalancer and consumer loses the ownership of the current partition. This may lead to duplicate processing when the partition is reassigned to consumer. There are some other operations such as database connection operation, file operation, or caching operations that may be part of consumer; you may want to deal with this before ownership of the partition is lost.

Kafka provides you with an API to handle such scenarios. It provides the `ConsumerRebalanceListener` interface that contains the `onPartitionsRevoked()` and `onPartitionsAssigned()` methods. We can implement these two methods and pass an object while subscribing to the topic using the `subscribe` method discussed earlier:

```

import org.apache.kafka.clients.consumer.ConsumerRebalanceListener;
import org.apache.kafka.common.TopicPartition;

import java.util.Collection;

public class DemoRebalancer implements ConsumerRebalanceListener {
    @Override
    public void onPartitionsRevoked(Collection<TopicPartition> collection) {
        //TODO: Things to Do before your partition got revoked
    }

    @Override
    public void onPartitionsAssigned(Collection<TopicPartition> collection) {
        //TODO : Things to do when new partition get assigned
    }
}

```

Summary

This concludes our section on Kafka producers and consumers. This lab addresses one of the key functionalities of Kafka message flows.

In the next lab, we will go through an introduction to Spark and Spark streaming, and then we will look at how Kafka can be used with Spark for a real-time use case and the different ways to integrate Spark with Kafka.