

Lab 7-1. Using Kafka with Confluent Platform



Use this quick start to get up and running with Confluent Platform and its main components in a development environment. This quick start uses Confluent Control Center included in Confluent Platform for topic management and event stream processing using ksqlDB.

In this quick start, you create Apache Kafka® topics, use Kafka Connect to generate mock data to those topics, and create ksqlDB streaming queries on those topics. You then go to Control Center to monitor and analyze the streaming queries.

Step 1: Start Confluent Platform

1. Open new terminal and go to following directory:

```
cd /headless/kafka-advanced/confluent-6.1.1/bin
```

2. Set the environment variable for the Confluent Platform directory.

```
export CONFLUENT_HOME=/headless/kafka-advanced/confluent-6.1.1
```

3. Install the Kafka Connect Datagen source connector using the Confluent Hub client. This connector generates mock data for demonstration purposes and is not suitable for production. ConfluentHub is an online library of pre-packaged and ready-to-install extensions or add-ons for Confluent Platform and Kafka.

```
./confluent-hub install --no-prompt confluentinc/kafka-connect-datagen:latest
```

4. Start Confluent Platform using the Confluent CLI `confluent local services start` command. This command starts all of the Confluent Platform components, including Kafka, ZooKeeper, Schema Registry, HTTP REST Proxy for Kafka, Kafka Connect, ksqlDB, and Control Center.

```
./confluent local services start
```

Run above command again if any of the following services fail to start first time.

Your output should resemble:

```
Starting Zookeeper
Zookeeper is [UP]
Starting Kafka
Kafka is [UP]
Starting Schema Registry
Schema Registry is [UP]
Starting Kafka REST
Kafka REST is [UP]
Starting Connect
Connect is [UP]
Starting KSQL Server
KSQL Server is [UP]
Starting Control Center
Control Center is [UP]
```

Step 2: Create Kafka Topics

In this step, you create Kafka topics using Confluent Control Center. Confluent Control Center provides the functionality for building and monitoring production data pipelines and event streaming applications.

1. Navigate to the Control Center web interface at <http://localhost:9021>.

If you installed Confluent Platform on a different host, replace `localhost` with the host name in the address.

It may take a minute or two for Control Center to come online.

Note

Control Center won't connect to ksqlDB if Control Center isn't open and running in a `localhost` browser session.

2. Click the **controlcenter.cluster** tile.

Home

1

Healthy clusters

0

Unhealthy clusters

 Search cluster name or id

controlcenter.cluster

Running

Overview

Brokers

1

Partitions

321

Topics

61

Production

17.53KB/s

Consumption

13.57KB/s

Connected services

ksqlDB clusters

1

Connect clusters

1

3. In the navigation bar, click **Topics** to open the topics list, and then click **Add a topic**.

Cluster overview
Brokers
Topics
Connect
ksqlDB
Consumers
Replicators
Cluster settings

Topics

☒ Hide internal topics

+ Add a topic

Topics Topic name	Availability			Throughput			Consumed
	Partitions	Under replicated	Followers	Out of sync	Observers	Out of sync	
<u>default ksql_processing_log</u>	1	0	1	0	0	0	--

4. In the Topic name field, specify `pageviews` and click **Create with defaults**.

Note that topic names are case-sensitive.

CONTROLCENTER.CLUS...

Overview

Brokers

Topics

Connect

ksqldb

Consumers

Replicators

Cluster settings

ALL TOPICS >

New topic

Topic name*

Number of partitions*

Create with defaults Customize settings [Cancel](#)

5. In the navigation bar, click **Topics** to open the topics list, and then click **Add a topic**.

6. In the Topic name field, specify `users` and click **Create with defaults**.

Step 3: Install a Kafka Connector and Generate Sample Data

In this step, you use Kafka Connect to run a demo source connector called `kafka-connect-datagen` that creates sample data for the Kafka topics `pageviews` and `users`.

1. Run the first instance of the Kafka Connect Datagen connector to produce Kafka data to the `pageviews` topic in AVRO format.
 1. In the navigation bar, click **Connect**.
 2. Click the `connect-default` cluster in the **Connect Clusters** list.
 3. Click **Add connector**.
 4. Select the `DatagenConnector` tile.
 5. In the **Name** field, enter `datagen-pageviews` as the name of the connector.
 6. Enter the following configuration values: (Scroll down browser screen to view all fields and also leave other fields empty)
 - **Key converter class:** `org.apache.kafka.connect.storage.StringConverter`.
 - **kafka.topic:** `pageviews`.
 - **max.interval:** `100`.

- **quickstart:** `pageviews` .

Add Connector

01 Setup connection

02 Test and verify

How should we connect to your data?

Connector class ⓘ

`io.confluent.kafka.connect.datagen.DatagenConnector` ⓘ

Name

`datagen-pageviews`

Common

Tasks max ⓘ

Key converter class ⓘ

`org.apache.kafka.connect.storage.StringConverter`

Value converter class ⓘ

Header converter class ⓘ

How should we connect to your data?

Common

Transforms

Predicates

Error Handling

Topic Creation

General

Additional Properties

7. Click **Continue**.

8. Review the connector configuration and click **Launch**.

CONNECT CLUSTERS > CONNECT-DEFAULT > CONNECTORS > SOURCES >

Add Connector

01 SETUP CONNECTION

02 TEST AND VERIFY

```
{
  "name": "datagen-pageviews",
  "connector.class": "io.confluent.kafka.connect.datagen.DatagenConnector",
  "key.converter": "org.apache.kafka.connect.storage.StringConverter",
  "kafka.topic": "pageviews",
  "max.interval": "100",
  "iterations": "1000000000",
  "quickstart": "pageviews"
}
```

Launch

Back

[Download connector config file](#)

2. Run the second instance of the [Kafka Connect Datagen] connector to produce Kafka data to the `users` topic in AVRO format.

1. In the navigation bar, click **Connect**.
2. Click the `connect-default` cluster in the **Connect Clusters** list.
3. Click **Add connector**.
4. Select the `DatagenConnector` tile.
5. In the **Name** field, enter `datagen-users` as the name of the connector.
6. Enter the following configuration values:
 - **Key converter class:** `org.apache.kafka.connect.storage.StringConverter`
 - **kafka.topic:** `users`
 - **max.interval:** `1000`
 - **quickstart:** `users`
7. Click **Continue**.
8. Review the connector configuration and click **Launch** (Scroll at bottom of screen).

Step 4: Create and Write to a Stream and Table using ksqlDB

Tip

You can also run these commands using the ksqlDB CLI from your terminal with this command:

```
/headless/kafka-advanced/confluent-6.1.1/bin/ksql http://localhost:8088 .
```

Create Streams and Tables

In this step, you use ksqlDB to create a stream for the `pageviews` topic and a table for the `users` topic.

1. In the navigation bar, click **ksqlDB**.
2. Select the `ksqlDB` application.
3. Copy the following code into the editor window and click **Run query** to create the `PAGEVIEWS` stream. Stream names are not case-sensitive.

```
CREATE STREAM PAGEVIEWS
  (VIEWTIME BIGINT, USERID VARCHAR, PAGEID varchar)
  WITH (KAFKA_TOPIC='pageviews', VALUE_FORMAT='AVRO');
```

4. Copy the following code into the editor window and click **Run query** to create the `USERS` table. Table names are not case-sensitive.

```
CREATE TABLE USERS
  (USERID VARCHAR PRIMARY KEY, REGISTERTIME BIGINT, GENDER VARCHAR, REGIONID
  VARCHAR)
  WITH (KAFKA_TOPIC='users', VALUE_FORMAT='AVRO');
```

Write Queries

In this step, you create ksqldb queries against the stream and the table you created above.

1. In the **Editor** tab, click **Add query properties** to add a custom query property.
2. Set the `auto.offset.reset` parameter to `Earliest`.

The setting instructs ksqldb queries to read all available topic data from the beginning. This configuration is used for each subsequent query.



3. Create the following queries.

1. Click **Stop** to stop the current running query.
2. Create a non-persistent query that returns data from a stream with the results limited to a maximum of three rows:

Enter the following query in the editor:

```
`SELECT PAGEID FROM PAGEVIEWS EMIT CHANGES LIMIT 3;`
```

3. Click **Run query**. Your output should resemble:

Data structure STREAM	 <input type="text" value="Filter by keyword"/>	
Total messages --	PAGEID	
Messages/sec --	Page_71	
Total message bytes --	Page_13	
Message fields • PAGEID	Page_52	

Click the **Card view** or **Table view** icon to change the output layout.

4. Create a persistent query (as a stream) that filters the

```
`PAGEVIEWS` stream for female  
users. The results from this query are written to the Kafka  
`PAGEVIEWS_FEMALE` topic:
```

Enter the following query in the editor:

```
CREATE STREAM PAGEVIEWS_FEMALE  
AS SELECT USERS.USERID AS USERID, PAGEID, REGIONID  
FROM PAGEVIEWS LEFT JOIN USERS ON PAGEVIEWS.USERID = USERS.USERID
```

```
WHERE GENDER = 'FEMALE'
EMIT CHANGES;
```

5. Click **Run query**. Your output should resemble:

```
0 | {
1 |   "@type": "currentStatus",
2 |   "statementText": "CREATE STREAM PAGEVIEWS_FEMALE WITH (KAFKA_TOPIC='PAGEVIEWS_
3 |   "commandId": "stream/'PAGEVIEWS_FEMALE'/create",
4 |   "commandStatus": {
5 |     "status": "SUCCESS",
6 |     "message": "Created query with ID CSAS_PAGEVIEWS_FEMALE_0"
7 |   },
8 |   "commandSequenceNumber": 6,
9 |   "warnings": []
10 | }
```

6. Create a persistent query where `REGIONID` ends with `8` or `9`. Results from this query are written to the Kafka topic named `pageviews_enriched_r8_r9` as explicitly specified in the query:

Enter the following query in the editor:

```
CREATE STREAM PAGEVIEWS_FEMALE_LIKE_89
  WITH (kafka_topic='pageviews_enriched_r8_r9', value_format='AVRO')
  AS SELECT * FROM PAGEVIEWS_FEMALE
  WHERE REGIONID LIKE '%_8' OR REGIONID LIKE '%_9'
  EMIT CHANGES;
```

7. Click **Run query**. Your output should resemble:

```
0 | {
1 |   "@type": "currentStatus",
2 |   "statementText": "CREATE STREAM PAGEVIEWS_FEMALE_LIKE_89 WITH (KAFKA_TOPIC='pag
3 |   "commandId": "stream/'PAGEVIEWS_FEMALE_LIKE_89'/create",
4 |   "commandStatus": {
5 |     "status": "SUCCESS",
6 |     "message": "Created query with ID CSAS_PAGEVIEWS_FEMALE_LIKE_89_7"
7 |   },
8 |   "commandSequenceNumber": 8,
9 |   "warnings": []
10 | }
```

8. Create a persistent query that counts the `PAGEVIEWS` for each `REGION` and `GENDER`

combination in a tumbling window of 30 seconds when the count is greater than 1. Because the procedure is grouping and counting, the result is now a table, rather than a stream. Results from this query are written to a Kafka topic called `PAGEVIEWS_REGIONS`:

Enter the following query in the editor:

```
CREATE TABLE PAGEVIEWS_REGIONS
AS SELECT GENDER, REGIONID , COUNT(*) AS NUMUSERS
FROM PAGEVIEWS LEFT JOIN USERS ON PAGEVIEWS.USERID = USERS.USERID
WINDOW TUMBLING (size 30 second)
GROUP BY GENDER, REGIONID
HAVING COUNT(*) > 1
EMIT CHANGES;
```

9. Click **Run query**. Your output should resemble:

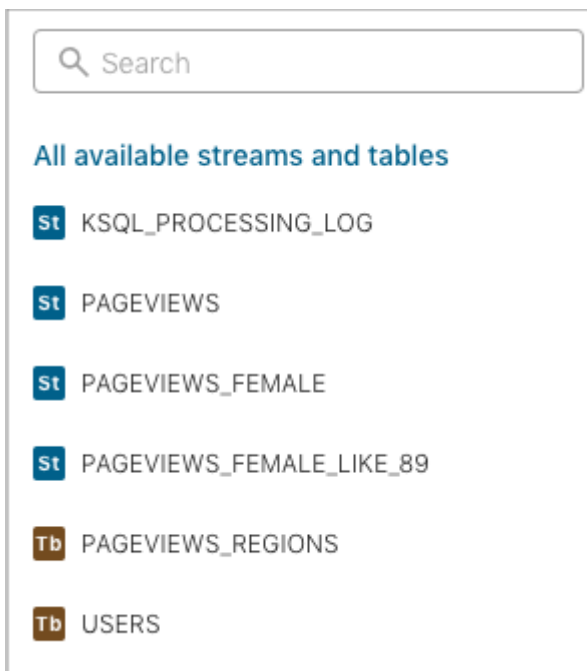
```
0  {
1    "@type": "currentStatus",
2    "statementText": "CREATE TABLE PAGEVIEWS_REGIONS WITH (KAFKA_TOPIC='PAGEVIEWS_F
3    "commandId": "table/'PAGEVIEWS_REGIONS'/create",
4    "commandStatus": {
5      "status": "SUCCESS",
6      "message": "Created query with ID CTAS_PAGEVIEWS_REGIONS_9"
7    },
8    "commandSequenceNumber": 10,
9    "warnings": []
10 }
```

10. Click the **Running queries** tab. You should see the following persisted queries:

- PAGEVIEWS_FEMALE
- PAGEVIEWS_FEMALE_LIKE_89
- PAGEVIEWS_REGIONS

11. Click the **Editor** tab. The **All available streams and**

tables pane shows all of the streams and tables that you can access.



12. In the **All available streams and tables** section, click **KSQL_PROCESSING_LOG** to view the stream's schema, including nested data structures.

Run Queries

In this step, you run the ksqlDB queries you save as streams and tables above in the previous section.

1. In the **Streams** tab, select the `PAGEVIEWS_FEMALE` stream.
2. Click **Query stream**.

The editor opens, and streaming output of the query displays.

3. Click **Stop** to stop the output generation.
4. In the **Tables** tab, select `PAGEVIEWS_REGIONS` table.
5. Click **Query table**.

The editor opens, and streaming output of the query displays.

6. Click **Stop** to stop the output generation.

Step 5: Monitor Consumer Lag

1. In the navigation bar, click **Consumers** to view the consumers created by ksqlDB.
2. Click the consumer group ID to view details for the `_confluent-ksql-default_query_CSAS_PAGEVIEWS_FEMALE_5` consumer group.

From the page, you can see the consumer lag and consumption values for your streaming query.